

Dr. Alexandra Chen

Email: a.chen@stanford.edu • **Phone:** +1 (650) 555-0142 • **Location:** Stanford, CA

Web: alexchen.io • **GitHub:** github.com/alexchen • **LinkedIn:** linkedin.com/in/alexchen • **Google Scholar**

Summary

Machine learning researcher with 8+ years of experience in natural language processing and large language models. Published 25+ papers at top venues (NeurIPS, ICML, ACL, EMNLP) with 3,800+ citations. Led research teams developing foundation models deployed to millions of users. Seeking faculty positions in computer science.

Education

Ph.D. in Computer Science

Stanford University

2016 – 2021

Stanford, CA

- Thesis: “Efficient Transformers for Long-Range Sequence Modeling”
- Advisor: Prof. Christopher Manning Committee: Profs. Percy Liang, Dan Jurafsky
- Stanford Graduate Fellowship (SGF) recipient

M.S. in Computer Science

Carnegie Mellon University

2014 – 2016

Pittsburgh, PA

- Focus: Machine Learning and Language Technologies
- GPA: 4.0/4.0 Dean’s List all semesters

B.S. in Mathematics and Computer Science

MIT

2010 – 2014

Cambridge, MA

- Summa Cum Laude Phi Beta Kappa
- Undergraduate thesis on spectral graph theory (awarded departmental prize)

Professional Experience

Senior Research Scientist

Google DeepMind

2023 – Present

Mountain View, CA

- Lead a team of 8 researchers developing next-generation language models
- Designed novel attention mechanisms reducing compute by 40% at scale
- Core contributor to Gemini model family; responsible for efficiency improvements
- Mentored 4 PhD interns, all placed at top research labs or faculty positions

Research Scientist

Meta AI (FAIR)

2021 – 2023

Menlo Park, CA

- Developed sparse mixture-of-experts models for efficient inference
- Led the retrieval-augmented generation (RAG) project for factual consistency
- Published 8 papers; 2 received spotlight/oral designations at NeurIPS
- Open-sourced model checkpoints downloaded 500K+ times

Research Intern

OpenAI

Summers 2018, 2019

San Francisco, CA

- Contributed to GPT-2 training infrastructure and evaluation framework
- Developed techniques for reducing hallucination in text generation

Selected Publications

1. **A. Chen**, J. Smith, R. Patel. “FlashLinear: Sub-Quadratic Attention via Structured State Spaces.” *NeurIPS 2024*. (**Oral, top 0.5%**)
2. **A. Chen**, M. Johnson. “Scaling Laws for Mixture-of-Experts Language Models.” *ICML 2024*.
3. L. Wang, **A. Chen**, K. Lee. “Retrieval-Augmented Generation with Learned Relevance.” *ACL 2023*. (**Best Paper Award**)
4. **A. Chen**, T. Brown, S. Narang. “Efficient Long-Context Transformers via Sliding Window Attention.” *NeurIPS 2022*. (**Spotlight**)
5. **A. Chen**, C. Manning. “Linear Transformers with Gated Retention.” *ICLR 2021*.

Awards & Honors

- ACL Best Paper Award, 2023
- MIT Technology Review 35 Under 35, 2024
- Stanford Graduate Fellowship, 2016–2021
- NSF Graduate Research Fellowship, 2016
- CRA Outstanding Undergraduate Award, 2014
- Phi Beta Kappa, MIT Chapter, 2014

Technical Skills

Languages: Python, C++, CUDA, Julia, Rust
Frameworks: PyTorch, JAX, TensorFlow, Hugging Face, DeepSpeed, Megatron-LM
Infrastructure: Kubernetes, Slurm, TPU Pods, multi-node GPU training (up to 2048 GPUs)
Specialties: Large language models, efficient attention, distributed training, NLP

Academic Service

- **Area Chair:** NeurIPS 2024, ICML 2024, ACL 2023
- **Reviewer:** NeurIPS, ICML, ICLR, ACL, EMNLP, JMLR, TACL (100+ reviews total)
- **Workshop Organizer:** “Efficient NLP” at EMNLP 2023; “Scaling Laws” at ICML 2024
- **Invited Talks:** Google, Microsoft, Meta, Amazon, Apple, NVIDIA, 15+ universities